

Green Shift Left - Evaluating energy efficiency on the web with static analysis

Jort van Driel
TU Delft
jortdriel@tudelft.nl

Dorian Erhan
TU Delft
derhan@tudelft.nl

Weicheng Hu
TU Delft
weichenghu@tudelft.nl

Giannos Rekkas
TU Delft
grekkas@tudelft.nl

Abstract—We asked for efficiency, but the web grew bloated. We asked for simplicity, but JavaScript expanded into a sprawling, inefficient beast. The result? Web applications that consume more energy than necessity, contributing significantly to global greenhouse gas emissions. So, we asked: Why is there no tool that helps developers understand the energy impact of their code? And then we set out to create one. This paper presents a deep dive into JavaScript coding patterns and their energy footprints, showing how minor adjustments can yield significant energy savings. We tested these patterns, compared them, and built an ESLint-based static analysis tool to flag inefficient code in real time. In our experiments, we found that replacing GIFs with more efficient formats and using document fragments can drastically reduce energy consumption. Our work aims to arm developers with the insights needed to make energy-efficient decisions, driving a more sustainable web without sacrificing performance.

Index Terms—JavaScript energy efficiency, Web optimization, Static analysis, ESLint, Sustainable software development, Coding patterns, Web performance

I. INTRODUCTION

Each year, hardware advances bring faster processors and larger memory capacities, yet no matter how much improvement has been achieved on hardware, developers always have the tendency to add functionality to make their software push the performance boundaries [1]. This has led to increasingly bloated software, where efficiency and optimization often take a back seat to rapid development and feature richness. This shift has resulted in an industry-wide lack of concern for efficiency and simplicity, where there is still an ever-increasing demand for performance optimization despite the employment of faster CPUs and larger memory systems. In stark contrast, the Apollo Guidance Computer successfully sent humans to the moon with just 4KB of RAM, while today a single browser tab can consume gigabytes of memory without much scrutiny.

Although the Web is a fundamental part of modern software development, research on its energy efficiency patterns remains scarce [2]. Web applications increasingly rely on excessive JavaScript frameworks, heavy multimedia content, and inefficient rendering techniques, all of which require significant processing power. As a result, the Internet’s energy consumption has skyrocketed and now accounts for approximately 3.7% of global greenhouse gas emissions, equivalent to all of the world’s air traffic [3]. Given the scale of its environmental impact, optimizing the Web’s energy efficiency should be a top priority.

At the heart of this challenge is JavaScript, which has become the cornerstone of modern software development. From front-end interfaces to back-end services to mobile apps, the language is ubiquitous across the web development stack. According to Stack Overflow’s 2024 Developer Survey [4], JavaScript is the most widely used programming language for the tenth year in a row. Its ubiquity means that even small inefficiencies in JS code can add up to significant energy costs across millions of devices and users.

JavaScript is a multi-paradigm language that supports object-oriented, imperative, and declarative programming styles. Besides the fact that it is an interpreted language, another contributor to performance issues in JavaScript is its rich and often functionally redundant APIs. Multiple coding constructs can provide the same functionality, but differ significantly in performance [5]. Developers often use sub-optimal idioms, inadvertently introducing energy inefficiencies.

Inefficient software directly contributes to higher energy consumption. While software does not independently emit CO₂, it does drive hardware resource usage [6]. An analysis of the COP28 climate conference website found that bloated web code, such as unused scripts, oversized images, and heavy third-party libraries, caused excessive data transfer and processing [7].

Despite growing awareness of software sustainability, developers still lack feedback on the energy impact of their code. Existing tools, such as linters and static analyzers, focus on style, correctness, and maintainability, with little to no attention paid to energy efficiency. As a result, energy-related anti-patterns go undetected, and developers may be unaware that certain idioms or practices could have a measurable environmental impact.

This leaves a huge gap: the lack of energy awareness in these everyday developer tools like linters. Many developers may use static code analysis to determine which parts of their code can be optimized in terms of time complexity and concurrency. Although time complexity can indeed serve as a predictor of energy consumption in some cases [8]. However, there is no existing literature that suggests this correlation universally applies. Research indicates that energy consumption is influenced by factors beyond algorithmic complexity, such as hardware architecture, compiler optimizations, and system-level interactions [9], [10]. For example, the energy efficiency of an algorithm can vary depending on how it utilizes the

memory hierarchy and manages data movement [11], so we cannot conclude that time complexity is proportional to energy consumption.

This paper addresses the lack of energy-aware development practices in the Web ecosystem. First, to address the fact that there are currently few industry guidelines that systematically address energy-aware coding practices, we present an empirical analysis of the energy footprint associated with various JavaScript coding patterns and identify opportunities where alternative idioms yield meaningful energy savings.

Our second goal is to provide a low-overhead, open-source static analysis solution that integrates into existing CI/CD pipelines and gives developers immediate, actionable insight into energy-inefficient code segments.

Key Take-aways. The rapid expansion of the Internet has led to increasingly inefficient software development, where optimization is often overlooked in favor of feature-rich applications. This inefficiency is particularly evident in the modern Web, where excessive JavaScript frameworks, heavy multimedia content, and poor rendering techniques significantly contribute to high energy consumption, now responsible for approximately 3.7% of global greenhouse gas emissions. Given JavaScript’s ubiquity across the web ecosystem, even minor inefficiencies can scale into substantial environmental costs. Despite growing awareness of sustainable software practices, existing development tools lack mechanisms to provide real-time feedback on the energy impact of code. To address this gap, this research presents an empirical analysis of JavaScript coding patterns and proposes a low-overhead, open-source static analysis solution that integrates with CI/CD pipelines. By leveraging established energy-aware design patterns and providing actionable insights, this work aims to empower developers to write more efficient, sustainable code, ultimately reducing the environmental footprint of web applications.

All of the code used in this study, including the linter plugin, the experiments, the testing framework, and the results are open source and available to the public through a GitHub repository[†].

II. BACKGROUND AND RELATED WORK

The escalating energy consumption of web applications has become a significant concern within the software engineering community. As digital services proliferate, their environmental impact intensifies, leading to a rise in the need for the adoption of energy-efficient development practices. This section explores existing research and initiatives aimed at promoting sustainability in web development, with a particular focus on energy patterns, responsible design collectives, and static analysis tools.

Energy patterns are reusable solutions designed to reduce the energy footprint of software applications. While initially conceptualized for mobile applications, recent studies have investigated their applicability to web development. Rani et al. conducted an exploratory study to assess the transferability of

mobile energy patterns to the web domain. They identified 20 patterns that could be adapted and validated through interviews with six expert web developers. The study revealed that developers were particularly concerned with functional anti-patterns and emphasized the need for guidelines to detect such patterns in source code. Empirical evaluations of patterns like ‘Dynamic Retry Delay’ (DRD) and ‘Open Only When Necessary’ (OOWN) indicated that while OOWN led to energy savings, DRD did not exhibit a significant reduction in energy consumption [2].

The Collectif Conception Numérique Responsable (CNumR) is a French collective comprising experts and organizations dedicated to promoting responsible digital design. CNumR offers tools, certifications, and training to encourage eco-friendly web development practices. Their work includes the development of best practices and guidelines for reducing the environmental impact of digital services [12].

To bridge the gap between awareness and implementation of energy-efficient coding practices, we propose to use ESLint, a widely used static code analysis tool for JavaScript, to identify energy hotspots within codebases. ESLint works by constructing an abstract syntax tree from the source code, allowing us to effectively traverse and analyze the code structure. By developing custom ESLint rules that focus on patterns known to impact energy consumption, we hope to provide developers with actionable insights to help them make trade-offs between energy efficiency and other factors. A similar solution has also been proposed by the Green Code Initiative [13]. For our rules, we can also draw on many energy design patterns relevant to web development, such as those proposed by the Collectif Conception Numérique Responsable, also known as CNUMR [14].

Despite these advancements, challenges persist in integrating energy-efficient practices into mainstream web development. There is a notable lack of standardized guidelines and tools that seamlessly integrate into developers’ workflows, providing real-time feedback on the energy implications of their coding choices. Moreover, the effectiveness of existing energy patterns and tools varies, underscoring the need for continuous evaluation and refinement.

Building upon the existing work, our research aims to address these gaps by developing a low-overhead, open-source static analysis solution that integrates into existing CI/CD pipelines. This tool will provide developers with immediate, actionable insights into energy-inefficient code segments, facilitating energy-aware development practices in the web ecosystem.

III. METHODOLOGY

In this section, we justify why we chose ESLint as a linter to build our own rules, and how we set up our experiment and result processing.

A. Language and Linter Selection

An academic study measuring 27 languages found that performance and energy efficiency are not always aligned —

[†]<https://github.com/JortvD/cs4575-g5-p2>

a "faster" language is not necessarily the most energy efficient [15]. So, the choice of programming languages is not limited. **JavaScript** was chosen as the focus for this experiment. One of the motivations is that JavaScript is the most widely used of all programming languages [4].

Therefore, as the most popular linter for JavaScript, ESLint is used more often compared to other linters. Also, unlike some other languages, like for example, Python, where we would have to manually implement an AST with Pylint, ESLint already provides this functionality in its infrastructure, which simplifies our efforts. All of this leads to a better ecosystem for extending ESLint. We also write basic unit tests for the linting rules to make sure they work as desired before running the experiments.

B. Selected Design Patterns

In our research, we found several suggested design patterns for web design as well as a number of static code analysis rules that have already been implemented [13]. Here are the antipatterns and design patterns we decided to implement analysis rules for, and our motivation for choosing them. We also explain how we set up similar code samples that either comply or violate these rules to find the difference in energy consumption between the two scenarios.

- **Use lazy attribute:** We implement a rule that warns when using a `` element that is not set to load lazily. We found from the MDN web documentation that this can reduce energy consumption by delaying loading an image until it is visible [16]. Our rule-compliant experiment imports 16 images using `loading=lazy`, and our violating sample uses `loading=eager`, the latter being the default behavior.
- **Preload link:** Adding `rel=preload` to a `<link>` element for all documents you know will be loaded on your webpage was introduced as a form of speculative loading, and is said to improve performance [17]. We are interested in whether this performance increase has an impact on energy consumption, so we have two samples that import a stylesheet, one with a preload link and one without.
- **Avoid using canvas:** When loading a webpage, about half of the energy is used for rendering [18]. Developers use `<canvas>` in cases that require detailed graphical control or complex animations, but according to a previous study, CSS can have the best performance in terms of accuracy and precision among different web technologies [19]. We have set up a basic experiment comparing a canvas with a CSS animation to see what the effect is on energy consumption.
- **Avoid using GIF:** Animated GIFs were introduced in 1995, and there are more modern encodings available that are more compressed [14]. Inspired by design pattern 4002 from CNUMR, we experiment with how the energy consumption of the more modern `.webp` format compares to using `.gif`.
- **Use standard fonts:** Different font sizes impact client-side resource consumption, including CPU usage, memory allocation, load time, and energy consumption [20]. For this reason, we have implemented a rule that checks for custom fonts and issues a warning for fonts that are not in standard libraries. In our experiments, we either load a standard font or a custom font.
- **Avoid logging often:** `console.log()` is not standardized, so its behavior is browser-dependent. Depending on the browser engine, its functionality can be either synchronous or asynchronous. As JavaScript is single-threaded [21], multiple logging calls might end up having blocking behavior. This may lead to greater energy overhead compared to bundling multiple objects into a single logging operation.
- **Use document fragment:** When appending multiple elements to your HTML document using JavaScript, there are two options: appending directly to the document, and the lesser-known option of appending to a temporary document fragment that is then added to the document at once. The latter option can be more performant because it does not require the document to be refreshed many times [22], [23]. We create an experiment where we add 100,000 elements either directly to the document or via the document fragment.
- **Use intersection observer:** The Intersection Observer API can be used to create a listener for a part of the web page that is scrolled into view [24]. This allows you to load only when something is in view. The alternative solution is to check each time a user scrolls if something is in view with a scroll event listener. We will create two experiments to compare these two behaviors.
- **Use request animation frame:** The `requestAnimationFrame` functionality has been introduced to automatically update animations at the same rate as the display refresh rate. This API also introduces other changes, such as pausing when the page is in the background, which increases battery life [25]. We are testing to see if there is a difference in energy consumption when used in the foreground between `requestAnimationFrame` and the original `setInterval` behavior.
- **Avoid expensive identifiers:** A simple optimization is to cache calls that are expensive to make, such as finding items by some identifier using `querySelectorAll`. We create a simple experiment that toggles some items 500 million times, either calling `querySelectorAll` every time or only once before.
- **Avoid resizing images:** When an image is loaded into a web page, but then needs to be displayed at a different size, the image must be resized, causing some performance overhead. This is especially true when many images are displayed [14]. CNUMR's design pattern 34 recommends that images not be resized using their `width` and `height` attributes. We create an experiment to see if resizing versus not resizing has a significant

impact on energy consumption.

- **Respect the bfcache:** Finally, the back/forward cache allows for faster navigation when returning to a previously visited site by caching the resources and heap of that page. However, if an `unload` or `beforeunload` event listener is set, the bfcache will not be used, negating this performance benefit [26]. We test how this affects the energy consumption by testing it on a web page with a complex canvas scene.

C. Experimental Setup

The experiments were performed on an HP ZBook Studio G5 118L5ES running Ubuntu 24.04. This laptop has an Intel i9-9750H CPU, an Nvidia Quadro P2000 GPU, and 16 GB of RAM. To minimize background noise while running the experiments, we make sure that configuration settings are kept consistent across all experiment runs. When running the experiments, we make sure that the laptop is in a controlled environment, which includes

- No unnecessary applications/services running
- All external hardware disconnected from the device
- Notifications turned off
- Both wifi and Bluetooth were turned off
- Screen brightness display at 100%

D. Experimental Procedure

The goal is to measure the difference in energy consumption between inefficient JavaScript coding patterns and their supposedly optimized counterparts. For each comparison, we test a piece of code that follows a known anti-pattern against a more efficient version based on established coding best practices. For example, we take JavaScript code that loads GIFs and compare it to code that loads more efficient formats, such as WEBM. All measurements are performed using EnergiBridge [27], which is an open-source cross-platform energy measurement utility[†]. We chose joules as the unit of measurement because it provides an absolute amount of energy measurement that tells us exactly how much energy was used in total, independent of time variations.

For this study, we use Chromium version because it is the basis for many of the most widely used browsers, including Google Chrome, Microsoft Edge, and Opera, making it a good representative choice [28]. Testing on Chromium ensures that the results are relevant to a wide range of browsers built on the same engine. We use Chromium version 137.0.7104.0.

For each design pattern described in III-B, we perform an experimental test by launching Chromium and separately measuring energy consumption on an HTML page containing the anti-pattern and another HTML page containing the optimized version. To reduce bias, the order of the web pages is randomized during each trial. Because hardware temperature can affect energy consumption, the very first run (used as a warm-up) is excluded from the final analysis. In total, we repeated this step 30 times for each rule.

[†]<https://github.com/t Durieux/EnergiBridge>

Within each step, the following actions are performed:

- 1) Launch Chromium through EnergiBridge, getting 5 measurements per second.
- 2) Open a new tab containing the test site (either the anti-pattern or the optimized variant).
- 3) Wait 5 seconds to ensure that all resources are fully loaded.
- 4) Close the tab.
- 5) Wait 1 second for the tab to close properly.
- 6) Repeat steps 1-5 for the alternate variant (the one not tested in the previous step).
- 7) Close Chromium and wait 5 seconds before proceeding to the next step.

After each run, we reset all generated caches and user data before proceeding to the next run.

One of the rules involves testing the efficiency of code that is compatible and incompatible with the back-forward cache (bfcache). The bfcache is a browser optimization that stores pages in memory so that they do not have to be reloaded, allowing for instant backward and forward navigation [26]. Because of this, the methodology for this particular rule has an extra step where we navigate away and then back after loading the web page to make sure we are measuring this functionality correctly.

E. Output postprocessing

After running the experiments, we perform a number of data processing steps to obtain our results. First, we select only the samples from 1 second before the web page is opened in the browser, to ensure that we do not miss any relevant data because the web page may not be opened exactly at the millisecond we expect. The selection is made until the web page is closed, which is 11 seconds for all experiments except the bfcache experiment, where it is 17 seconds. We then sum the positive increases in `PACKAGE_ENERGY` (J) over this time period, which gives us an estimate of the CPU and measurable motherboard energy consumption per experiment. To remove any outliers, we use the Inter-Quartile Range(IQR) method to remove samples outside the upper and lower $1.5 * IQR$ whiskers. We manually check that this does not remove too many samples, but in our case, we did not accept any measurements that were rejected by the IQR calculation.

IV. RESULTS

In the table I, we show the results of the experiment. **All measurements are made in joules.** To avoid confusion, a violating rule corresponds to the anti-patterns described in III-B, while compliant rules represents the intended optimized version.

We calculate the Shapiro-Wilk test per experiment to see if the distribution of samples follows the normal distribution. We use the typical threshold of 0.05; any values above the threshold imply that the experiment is normally distributed. If that is the case, we apply Welch's two-sided t-test to see if the change in the mean between the two distributions is

significant, otherwise, there may still be a significant mean change that we can find with the two-sided Mann-Whitney U test.

Following the result, we provide an analysis of the percentage changes in joules before and after applying the rules in Table II. Here, in addition to the mean difference and the percentage mean change, we also apply Cohen's d to see if we can classify this change as large by including the variance of both distributions in the difference calculation.

Finally, we visualize Table I and II using bar charts to make a more straightforward comparison. The blue bar in Figure 1 represents the energy used if rules are violated, and the orange bar represents the energy used after rules are not violated. Figure 2 shows the percentage change in energy usage after the rules are complied with.

From Table I, we can see that apart from violating the animation frame, all other experiments are following a normal distribution, which means t-test p-values can be referred to in most cases. Here, rules that are statistically significant (p-value < 0.05) are summarized:

- **When energy consumption is higher if we comply with the rule:**
 - Avoid using canvas
 - Avoid using GIF
- **When energy consumption is higher if we violate the rule:**
 - Avoid expensive identifiers

Additionally, the request animation frame is not normally distributed. However, it has a U-test p-value of ≈ 0.001 , and its shape is close to a normally distributed graph according to its violin plot VII-A, so we see it as a weaker statement of statistical significance that suggests the energy consumption is lower if we compliant the rule.

Finally, we conduct a Cohen's d effect size analysis, the data can be found in Table II. An absolute effect size of ≈ 0.8 is considered small, anything bigger than 0.8 is considered a large effect [29].

- **Large Effect Sizes (absolute Cohen's d ≥ 0.8):**
 - **Lower Energy Consumption**
 - * Avoid expensive identifiers (Cohen's d = 25.115)
 - **Higher Energy Consumption**
 - * Avoid using canvas (Cohen's d = -1.019)
 - * Avoid using GIF (Cohen's d = -3.526)
- **Small Effect Sizes (absolute Cohen's d < 0.8):**
 - **Positive Effects (energy consumption decreases when applied):**
 - * Preload link (Cohen's d = 0.480)
 - * Use standard fonts (Cohen's d = 0.339)
 - * Avoid logging often (Cohen's d = 0.046)
 - * Respect the bfcache (Cohen's d = 0.25)
 - **Negative Effects (energy consumption increases when applied):**
 - * Use lazy attribute (Cohen's d = -0.354)

- * Use intersection observer (Cohen's d = -0.313)
- * Avoid resizing images (Cohen's d = -0.169)
- * Use document fragment (Cohen's d = -0.149)

From all the results, we can conclude that avoiding expensive identifiers leads to dramatically lower energy consumption. On the other hand, avoiding gifs and canvas leads to a higher energy consumption. The remaining rules do not seem to make a significant difference.

V. DISCUSSION AND LIMITATIONS

Key findings. Given Cohen's d effect size, is the improvement practically significant? Effect size analyses help assess practical significance, but might not be enough. Consider the context and explain in what sense the effect size might be relevant. E.g., to improve 2% in energy efficiency, the code will be less readable, or the user experience will not be so appealing. A particular method improves 2% but it will only be used 1% of the time.

Implications. This study has significant implications, particularly for developers working on web development, browser engines, or linters, as our findings contribute to improving coding practices, tooling efficiency, and overall software quality.

The results provide insights into energy-inefficient coding anti-patterns and their potential optimizations. These insights can guide updates to coding guidelines and industry standards, helping developers write cleaner, more energy-efficient code that better aligns with best practices. Additionally, our study also highlights that these idioms can be automatically detected and enforced by linters, reducing inefficiencies at scale. Understanding these patterns may also allow browser engines to mitigate the impact of energy-inefficient code or even let language developers substitute and deprecate the most problematic coding constructs.

Limitations. While our study provides insights into energy consumption on the web, we also need to acknowledge certain limitations that can affect the interpretation of the results.

First, the scope of the experiment is relatively small, mainly due to time constraints. Having only 12 rules is a good starting point for analyzing bad practices in web development, but it is not representative of the full range of energy-intensive behaviors on the web. It is also worth noting that the experiment was conducted in just one language (JavaScript), using one specific browser (Chromium) and one specific linter (ESLint). A broader investigation that includes different browser engines, frameworks, and linters would strengthen the generalizability of the results.

Second, another important consideration is that the experiments focus solely on measuring the energy consumption of JavaScript execution in the browser. Other energy-intensive factors, such as the cost of transmitting, decoding, and caching media assets like GIFs and WEBM, are not taken into account. Network-related energy consumption remains outside the scope of this study.

Thirdly, the examples used in the experiments are not representative of code found in production environments. The test instances consist of isolated patterns within an HTML page,

TABLE I
STATISTICAL RESULTS OF ENERGY CONSUMPTION (J) FOR COMPLAINT AND VIOLATING SAMPLES PER RULE

Experiment	Mean		Std Dev		Shapiro-Wilk		t-test	U-test
	Compliant	Violating	Compliant	Violating	Compliant	Violating	p	p
Use lazy attribute	11.624	11.494	0.330	0.400	0.760	0.223	0.194	
Preload link	11.102	11.240	0.298	0.276	0.454	0.561	0.087	
Avoid using canvas	11.448	11.081	0.276	0.430	0.560	0.554	< 0.001	
Avoid using GIF	44.546	42.167	0.655	0.694	0.572	0.489	< 0.001	
Use standard fonts	11.609	11.699	0.230	0.292	0.547	0.819	0.228	
Avoid logging often	11.344	11.360	0.422	0.247	0.126	0.610	0.867	
Use document fragment	77.685	77.332	2.473	2.274	0.387	0.14	0.566	
Use intersection observer	11.453	11.358	0.312	0.295	0.885	0.693	0.256	
Use request animation frame	35.657	37.548	0.834	2.062	0.63	0.033		< 0.001
Avoid expensive identifiers	96.843	160.024	1.696	3.127	0.880	0.625	< 0.001	
Avoid resizing images	11.691	11.641	0.250	0.331	0.623	0.297	0.533	
Respecting the BF-cache	289.987	291.696	6.926	6.756	0.487	0.856	0.337	

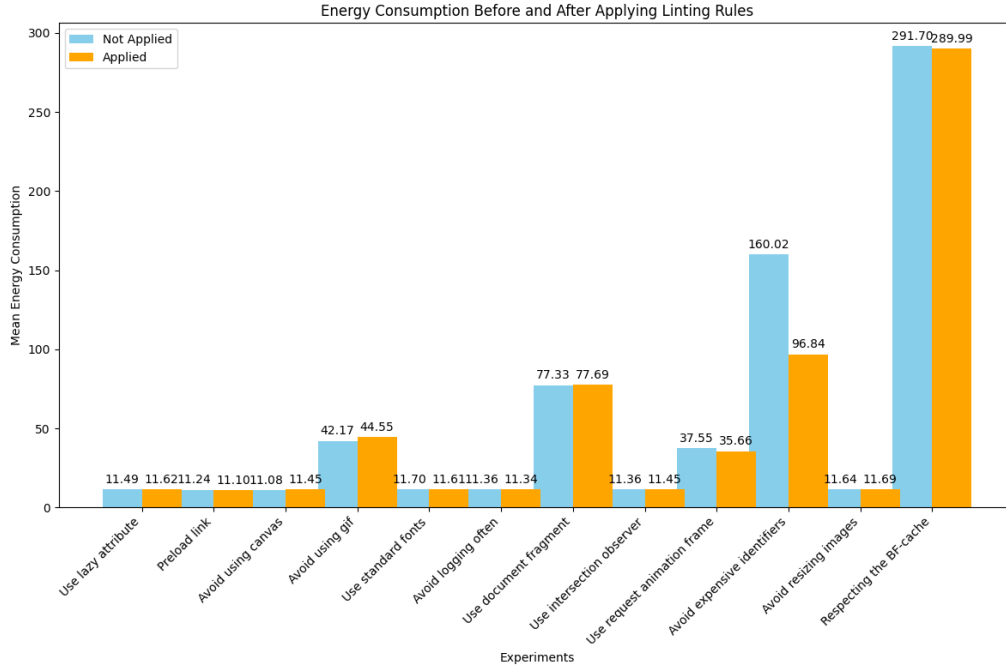


Fig. 1. Energy consumption comparison (joules) between the unoptimized and optimized rules

which may not fully reflect how these patterns are typically integrated into real-world websites. As a result, there may be differences in energy consumption when these elements interact with other components, scripts, or user behavior. To obtain more realistic results, future experiments should be conducted in real-world projects to evaluate the effectiveness of our findings in practical scenarios.

Furthermore, while the code used in these experiments is functional and provides meaningful insights, it is likely not optimal. There may be opportunities for improvements in efficiency and correctness within the codebase. Additionally, the experiment's implementation is relatively simple, primarily opening a browser and, in the case of the bfcache rule, navigating away and back. This functionality can be expanded

to account for more complex behavior, enabling a more comprehensive analysis of energy consumption in more diverse scenarios.

Finally, it is important to recognize that this work is only concerned with energy consumption; there is more to consider. For example, lazy loading of some web elements may indeed reduce energy consumption. However, according to a Cloudflare report, lazy loading can result in multiple server requests as users navigate the page, potentially adding overhead and impacting performance [30]. Therefore, the choice between implementation and energy consumption depends on the specific requirements of a project. It is also worth noting that there is a trade-off between efficiency/performance and maintainability. In some cases, the most energy-efficient




Fig. 2. Percentage change in energy consumption (joules)

TABLE II
STATISTICAL ANALYSIS OF ENERGY CONSUMPTION PER RULE

Experiment	Mean diff (ΔX)	Mean change (%)	Effect size (Cohen's d)
Use lazy attribute	-0.130	-1.117%	-0.354
Preload link	0.138	1.239%	0.480
Avoid using canvas	-0.368	-3.213%	-1.019
Avoid using GIF	-2.379	-5.34%	-3.526
Use standard fonts	0.089	0.767%	0.339
Avoid logging often	0.016	0.139%	0.046
Use document fragment	-0.354	-0.456%	-0.149
Use intersection observer	-0.095	-0.829%	-0.313
Use request animation frame	1.891	5.303%	N/A
Avoid expensive identifiers	63.180	65.240%	25.115
Avoid resizing images	-0.050	-0.424%	-0.169
Respect the bfcache	1.709	0.589%	0.25

solution may result in code that is harder to understand, modify, or debug. So, developers need to consider this for their codebases. Our goal is to educate users about potential energy hotspots so that they can make informed decisions based on their own needs.

VI. CONCLUSION AND FUTURE WORK

In conclusion, our work demonstrates that incorporating energy-awareness into standard web development practices can yield benefits in sustainable software engineering, thus reducing software’s environmental impact. By focusing on JavaScript and extending ESLint with custom rules, we identified how common coding patterns influence energy consumption in measurable ways. While some optimizations (such as using document fragments and caching expensive identifiers) led to striking improvements in overall energy usage, other suspected “green” patterns produced only minimal or even counterintuitive effects.

These findings highlight the complexity of energy-aware development: performance optimizations do not always translate

to lower energy consumption, and some best practices (e.g., avoiding GIFs entirely) may carry trade-offs in terms of feature set, user experience, or code maintainability. Although our experiments centered on only Chromium in controlled conditions, they illustrate the value of tool-supported, empirical investigations into software sustainability.

Future Work. Looking ahead, this research opens up several avenues for future work. The most valuable direction would be to integrate these experiments into real-world projects to validate the effectiveness of the approach in practical web applications rather than isolated test cases.

Similarly, extending the scope beyond JavaScript would be highly valuable. The vast majority of current websites rely on frameworks such as React or Angular. These frameworks introduce much more functionality and complexity because their patterns are independent of JavaScript. Thus, identifying energy hotspots in these tools can provide significant insight into optimizing real-world applications.

In addition, future studies could improve the experimental setup by running tests in more diverse environments, across different browsers (e.g., Firefox, Safari) and different devices, to assess the consistency of the results and mitigate potential biases. This would help determine whether energy-saving recommendations generalize across different platforms or require browser-specific optimizations.

Another natural extension would be to expand the set of rules evaluated to include a wider range of energy consumption patterns. Analyzing more complex functionality would provide a more comprehensive analysis of consumption across the web.

Finally, a user study could be conducted to assess whether developers find our linting tool useful and to get feedback on possible ways to improve the tool. Such a study would be crucial for understanding real-world usability, including whether the tool’s energy-saving rules introduce unintended trade-offs, such as reduced code readability, maintainability,

or even runtime performance. Understanding these factors is essential, as we hope this study encourages the push towards integrating energy-aware linting tools into existing continuous integration pipelines.

REFERENCES

- [1] H. Xu, N. Mitchell, M. Arnold, A. Rountev, and G. Sevitsky, “Software bloat analysis: Finding, removing, and preventing performance problems in modern large-scale object-oriented applications,” 11 2010, pp. 421–426.
- [2] P. Rani, J. Zellweger, V. Kousadianos, L. Cruz, T. Kehrler, and A. Bacchelli, “Energy patterns for web: An exploratory study,” in *Proceedings of the 46th International Conference on Software Engineering: Software Engineering in Society (ICSE-SEIS)*. IEEE, 2024, pp. 12–22. [Online]. Available: <https://arxiv.org/abs/2401.06482>
- [3] B. Couriol, “Co2.js helps developers track their application’s carbon footprint,” InfoQ, 05 2024. [Online]. Available: <https://www.infoq.com/news/2024/05/co2-js-carbon-footprint-release>
- [4] “Stack overflow developer survey,” 2024, accessed: 19 March 2025. [Online]. Available: <https://survey.stackoverflow.co/2024/technology#top-paying-technologies>
- [5] M. Selakovic and M. Pradel, “Performance issues and optimizations in javascript: An empirical study,” in *2016 IEEE/ACM 38th International Conference on Software Engineering (ICSE)*, 2016, pp. 61–72.
- [6] T. Imura, “10 recommendations for green software development,” *Green Software Foundation*, 11 2021. [Online]. Available: <https://greensoftware.foundation/articles/10-recommendations-for-green-software-development>
- [7] R. D. Caballar, “We need to decarbonize software: The way we write software has unappreciated environmental impacts,” *IEEE Spectrum*, vol. 61, no. 4, pp. 26–31, 2024.
- [8] K. Carter, S. M. G. Ho, M. M. A. Larsen, M. Sundman, and M. H. Kirkeby, “Energy and time complexity for sorting algorithms in java,” 2024. [Online]. Available: <https://arxiv.org/abs/2311.07298>
- [9] J. Pallister, S. J. Hollis, and J. Bennett, “Identifying compiler options to minimize energy consumption for embedded platforms,” *The Computer Journal*, vol. 58, no. 1, p. 95–109, Nov. 2013. [Online]. Available: <http://dx.doi.org/10.1093/comjnl/bxt129>
- [10] S. Hu and L. K. John, “Impact of virtual execution environments on processor energy consumption and hardware adaptation,” in *Proceedings of the 2nd International Conference on Virtual Execution Environments*, ser. VEE ’06. New York, NY, USA: Association for Computing Machinery, 2006, p. 100–110. [Online]. Available: <https://doi.org/10.1145/1134760.1134775>
- [11] T.-J. Yang, Y.-H. Chen, J. Emer, and V. Sze, “A method to estimate the energy consumption of deep neural networks,” in *2017 51st Asilomar Conference on Signals, Systems, and Computers*, 2017, pp. 1916–1920.
- [12] C. C. N. Responsable, “115 web ecodesign best practices,” 2024, accessed: 2025-04-03. [Online]. Available: <https://collectif.greenit.fr/ecoconception-web/>
- [13] Green-Code-Initiative, “Creedengo-javascript/eslint-plugin/readme.md at main · green-code-initiative/creedengo-javascript,” Mar 2025. [Online]. Available: <https://github.com/green-code-initiative/creedengo-javascript/blob/main/eslint-plugin/README.md>
- [14] CNUMR, *115 Web Ecodesign Best Practices*, GitHub, 2025, accessed: 2025-04-03. [Online]. Available: <https://github.com/cnumr/best-practices>
- [15] R. Pereira, M. Couto, F. Ribeiro, R. Rua, J. Cunha, J. a. P. Fernandes, and J. a. Saraiva, “Energy efficiency across programming languages: how do energy, time, and memory relate?” in *Proceedings of the 10th ACM SIGPLAN International Conference on Software Language Engineering*, ser. SLE 2017. New York, NY, USA: Association for Computing Machinery, 2017, p. 256–267. [Online]. Available: <https://doi.org/10.1145/3136014.3136031>
- [16] Mozilla, *Lazy Loading - Web Performance*, MDN Web Docs, 2025, accessed: 2025-04-03. [Online]. Available: https://developer.mozilla.org/en-US/docs/Web/Performance/Guides/Lazy_loading
- [17] —, *rel=preload - HTML: Hypertext Markup Language*, MDN Web Docs, Mar. 2025, accessed: 2025-04-03. [Online]. Available: <https://developer.mozilla.org/en-US/docs/Web/HTML/Attributes/rel/preload>
- [18] B. Poulain and S. Fraser, “How web content can affect power usage,” August 27 2019. [Online]. Available: <https://webkit.org/blog/8970/how-web-content-can-affect-power-usage/>
- [19] P. Garaizar, M. A. Vadillo, and D. L. de Ipiña, “Presentation accuracy of the web revisited: Animation methods in the HTML5 era,” *PLoS ONE*, vol. 9, no. 10, p. e109812, 2014. [Online]. Available: <https://doi.org/10.1371/journal.pone.0109812>

- [20] B. Dornauer, W. Vigl, and M. Felderer, “On the role of font formats in building efficient web applications,” 2023. [Online]. Available: <https://arxiv.org/abs/2310.06939>
- [21] MDN Web Docs contributors, *Main Thread*, Mozilla, Dec. 2024, accessed: 2025-04-03. [Online]. Available: https://developer.mozilla.org/en-US/docs/Glossary/Main_thread
- [22] Mozilla, *Document: createDocumentFragment() Method - Web APIs*, MDN Web Docs, Mar. 2024, accessed: 2025-04-03. [Online]. Available: <https://developer.mozilla.org/en-US/docs/Web/API/Document/createDocumentFragment>
- [23] MDN Web Docs contributors, *JavaScript Performance Optimization*, Mozilla, 2025, accessed: 2025-04-03. [Online]. Available: https://developer.mozilla.org/en-US/docs/Learn_web_development/Extensions/Performance/JavaScript#tips_for_writing_more_efficient_code
- [24] MDN Web Docs, *Intersection Observer API*, Mozilla, 2025, accessed: 2025-04-03. [Online]. Available: https://developer.mozilla.org/en-US/docs/Web/API/Intersection_Observer_API
- [25] MDN Web Docs contributors, *Window: requestAnimationFrame() method*, Mozilla, 2025, accessed: 2025-04-03. [Online]. Available: <https://developer.mozilla.org/en-US/docs/Web/API/Window/requestAnimationFrame>
- [26] —, *bfcache*, Mozilla, Sep. 2024, accessed: 2025-04-03. [Online]. Available: <https://developer.mozilla.org/en-US/docs/Glossary/bfcache>
- [27] J. Sallou, L. Cruz, and T. Durieux, “Energibridge: Empowering software sustainability through cross-platform energy measurement,” *arXiv preprint arXiv:2312.13897*, 2023. [Online]. Available: <https://arxiv.org/abs/2312.13897>
- [28] “What is chromium, and how does it enhance your browser?” Microsoft, 2025. [Online]. Available: <https://www.microsoft.com/en-us/edge/learning-center/what-is-chromium-how-does-it-enhance-your-browser?form=MA1312>
- [29] N. U. Library, “Cohen’s d - statistics resources,” accessed: April 3, 2025. [Online]. Available: <https://resources.nu.edu/statsresources/cohensd>
- [30] “What is lazy loading?” Cloudflare.com, 2024. [Online]. Available: <https://www.cloudflare.com/learning/performance/what-is-lazy-loading/>

VII. APPENDIX

A. Result violin plots

Each plot shows the results for the respective experiment. Note that for “all” each of the samples is included, while for “no outliers” we remove the outliers as described in section III-E.

